

Portfolio / Projects Overview


GovernOS, Agentic AI Planning & Orchestration System

Compiled natural-language goals into policy-gated DAG workflows across a multi-tenant agentic OS with schema validation, human approval gates, audit tracing, scoped runtime memory, and real-time replanning.

Inference Control Plane, LLM Infrastructure & Routing System

Developed a production-grade FastAPI gateway for LLM traffic control with priority-based routing, fallback routing, token optimization, Redis semantic caching, per-key rate limiting, and PostgreSQL audit logging.

Dashboard Screenshot



GuardrailX, Enterprise AI Governance Platform

Enforced runtime LLM governance across enterprise systems through prompt injection detection, jailbreak detection, PII redaction, content safety filtering, hallucination controls, and risk-based policy enforcement.

EnterpriseIQ, Enterprise Knowledge Intelligence Platform

Orchestrated LangGraph-based Agentic RAG using hybrid retrieval, cross-encoder reranking, local LLM inference, RBAC enforcement, and citation grounding across PDF, SQL, CSV, and JSON knowledge sources.

AI Infrastructure Hypervisor, GPU-Aware Virtualization Platform

Constructed a GPU-aware AI runtime using KVM/QEMU to provision isolated vLLM workloads with automated orchestration, Prometheus/Grafana observability, and workload scheduling across inference environments.